

Problem Set 3

1 MIP vs PCP

The goal of this question is to exhibit a close relationship between PCPs and a natural extension of interactive proofs in which the verifier interacts with multiple non-communicating provers.

Definition 1.1. A multi-prover interactive proof (MIP) for a language L consists of a probabilistic polynomial-time verifier V and k (computationally unbounded) provers P_1, \dots, P_k . Given a common input x , the verifier V is allowed to interact with each of the provers P_i . The provers are not allowed to interact with each other. We require that:

- **Completeness:** if $x \in L$ then V accepts with probability 1 when interacting with P_1, \dots, P_k .
- **Soundness:** if $x \notin L$, then for every P_1^*, \dots, P_k^* the verifier accepts with probability at most $1/2$ when interacting with P_1^*, \dots, P_k^* .

We say that an MIP is 2-message if it consists of a first message in which V sends queries to each of the k provers (simultaneously) and then a second message in which V gets back answers from each prover and then decides whether to accept or reject. We denote the length of messages sent from the verifier to each of the provers by ℓ_V and the length of messages sent back from the provers by ℓ_P .

Prove the following two claims:

1. If L has a k -prover 2-message MIP in which the verifier messages have length ℓ_V and the prover messages have length ℓ_P , then L has a $(k \cdot \ell_P)$ -query PCP where the PCP proof string has length $k \cdot 2^{\ell_V} \cdot \ell_P$.
2. If L has a q -query (non-adaptive) PCP in which the PCP has length at most m , then L has a $\text{poly}(q)$ -prover 2-message MIP in which the verifier messages have length $O(\log m)$ and prover messages have length 1 (i.e., a single bit).

Simplifying assumption: you may assume for simplicity that each of the PCP verifier's queries *by itself* is uniformly distributed (but the q queries may still be highly correlated).

Hint: have the provers answer according to the PCP but check consistency.

Guideline 1.2 (For Part 2).

1. First construct an MIP with q provers but very large soundness error $1 - \frac{1}{\text{poly}(q)}$ (later we will reduce the soundness error by repetition).
2. The q honest MIP provers answer according to the PCP.
3. With probability $1/2$, the MIP verifier simply runs the PCP verifier.

4. With probability $1/2$, the MIP verifier chooses a random coordinate $i \in [m]$, where m is the length of the PCP string, and sends i to all of the provers. It accepts iff all the provers answer with the same value.
5. To analyze soundness, fix $x \notin L$ and (wlog deterministic) cheating provers P_1^*, \dots, P_k^* . Analyze the following two cases separately:
 - (a) There exists a pair of provers that answer more than a $\frac{1}{100q^2}$ fraction of their queries inconsistently with one another.
 - (b) There does not exist such a pair of provers. To argue soundness in this case, first argue that on most coordinates all of the provers agree. Then, use this fact to construct an accepting PCP for x .
6. Finally, repeat the entire proof systems $\text{poly}(q)$ times (with $\text{poly}(q)$ sets of q provers) to obtain constant soundness error.

Remark 1.3. One can get rid of the simplifying assumption by choosing i in the consistency check differently - generate the q PCP queries by running the verifier and then output one of these q queries at random.

2 Tensor Codes

This exercise introduces and discusses the *tensoring* operation of codes. In a nutshell, this is a way to take two (linear) codes and combine them in an interesting way to derive a new code, inheriting some of the properties of the base codes and also magically generating some new properties.

Let \mathbb{F} be a finite field, and let $C_1 : \mathbb{F}^{k_1} \rightarrow \mathbb{F}^{n_1}$ and $C_2 : \mathbb{F}^{k_2} \rightarrow \mathbb{F}^{n_2}$ be linear codes. We construct a new code $C^\otimes : \mathbb{F}^{k_1 \times k_2} \rightarrow \mathbb{F}^{n_1 \times n_2}$ as follows: view the input as a $k_1 \times k_2$ sized matrix. Encode each one of the k_1 rows using C_2 to get a matrix of size $k_1 \times n_2$. Now encode each one of these n_2 columns to get the codeword (a matrix of size $n_1 \times n_2$). The resulting code is denoted by $C_1 \otimes C_2$.

Notation. Consider a codeword $c \in C_1 \otimes C_2$. Since we view c as a matrix, it will be convenient to index its entries accordingly. Thus, we use $c(i_1, i_2)$ to denote the (i_1, i_2) entry of this matrix.

Also, following the coding literature, if C is a code, we sometimes abuse notation and use C to refer both to the function (mapping messages to codewords) and to the corresponding set of codewords (i.e., the image of the function).

2.1 A Characterization

1. Consider the set \mathcal{M} of matrices $\mathbb{F}^{n_1 \times n_2}$ such that all of the rows are codewords of C_2 and all of the columns are codewords of C_1 . Show that $C_1 \otimes C_2 \subseteq \mathcal{M}$.

Hint: It may be useful to first show that the exact same code would be obtained if you first encoded the columns and then the rows.

2. Show that $\mathcal{M} \subseteq C_1 \otimes C_2$.

Hint: Recall that for linear codes, any linear combination of codewords is also a codeword.

3. Deduce that $\mathcal{M} = C_1 \otimes C_2$.

Hint: Don't forget your Matka!

2.2 Rate

Let $r_1 = k_1/n_1$ and $r_2 = k_2/n_2$ denote the rates of C_1 and C_2 , respectively. Show that C^\otimes is (1) a linear code and (2) has rate $r_1 \cdot r_2$.

2.3 Distance

Let $\delta_1, \delta_2 \in (0, 1)$ be the minimal relative distances of C_1 and C_2 , respectively. Show that C^\otimes has relative distance at least $\delta_1 \cdot \delta_2$.

Hint: Recall that for linear codes it suffices to bound the minimal number of zeros in a non-zero codeword.

In the following we use the notation $C^{\otimes d} = C \otimes C^{\otimes d-1}$, where $C^{\otimes 1} = C$. The code $C^{\otimes d}$ is sometimes called the d -th dimensional tensor of C . Note that if C has message length k and codeword length n then $C^{\otimes d}$ has message length k^d and codeword length n^d . Indeed, it is often more convenient to view the messages and codewords of $C^{\otimes d}$ as “ d -dimensional” matrices aka tensors (rather than viewing them as vectors),¹ In terms of notation, similarly to the case of $d = 2$ discussed above, for $c \in C^{\otimes d}$ and $i_1, \dots, i_d \in [n]$ we use $c(i_1, \dots, i_d)$ to denote the corresponding entry of c .

2.4 Tensor of Reed-Solomon: NOT TO BE SUBMITTED

Recall that the Reed-Solomon code $RS : \mathbb{F}^k \rightarrow \mathbb{F}^n$ consists of all degree $k-1$ univariate polynomials over \mathbb{F} . Describe in 10 words or less the image of the code $RS^{\otimes 2}$ (without referring to the tensoring operation of course).

2.5 Sumcheck for Tensor Codes

Let $C_{\text{base}} : \mathbb{F}^k \rightarrow \mathbb{F}^n$ be a code and consider the code $C : \mathbb{F}^{k^d} \rightarrow \mathbb{F}^{n^d}$ defined as $C = (C_{\text{base}})^{\otimes d}$. In this question we will show that C supports a “sumcheck-like” protocol - namely, a protocol for verifying the sum of some product-set of the codeword elements.

Let $S \subseteq [n]$. Construct an interactive proof for verifying the sum over the elements of $c \in C$ belonging to the set S^d : that is, an interactive proof for the statement $\sum_{i_1, \dots, i_d \in [S]} c(i_1, \dots, i_d) = \alpha$, for a given $\alpha \in \mathbb{F}$. That is, show an interactive proof in which the verifier has oracle access to $c \in C$ and explicitly given α and the prover is given both c and α . The verifier should accept if the claim is true and should reject with probability at least δ^d if the claim is false, no matter what the prover does (in all cases you may assume that c is indeed a valid codeword). The verifier's work should be $\text{poly}(d, n, \log(|\mathbb{F}|))$.

¹For $d > 3$, best not to try to visualize them ☹.

Hint: It may be useful to first think of the case $d = 2$, in which case a useful proof string to consider is $w \in \mathbb{F}^n$ defined as $w_\lambda = \sum_{i_2 \in S} c(\lambda, i_2)$, for every $\lambda \in [n]$. Handle larger dimensions by recursion.

Remark 2.1. *The detection probability δ^d may seem very low, but is actually fine as long as we choose the base code to have distance $\delta \gg 1 - 1/d$.*

Remark 2.2. *Note the remarkable fact that even if we start with a base code that does not support a “sumcheck-like” operation, the resulting tensor code magically does!*

Similarly, even if we start with a base code that is not locally-testable, the resulting tensor code² turns out to be locally testable.

2.6 Reflection: NOT TO BE SUBMITTED

Try to see how the sumcheck protocol that we saw for polynomials is a special case of the protocol for general tensor codes.

²Actually, this fails for $d = 2$, but works for $d \geq 3$.